Test Run Report

Test Run

RUN TIMESTAMP	JUDGE
2026-04-19T17:14:29.288199+00:00	gpt-5.4

PROMPTS	CANDIDATES
4	2

⚠ **Potential judge bias:** The judge (**gpt-5.4**) appears to share a model family with candidate(s) **GPT-5.4-mini, GPT-5.4**. LLMs are known to exhibit self-preference bias. Consider re-running with a different-family judge for an independent baseline.

EXECUTIVE SUMMARY

Global Score formula: 70% Judge Score (holistic quality) +

15% Cost Score (normalised inverse) +

15% Performance Score (latency + throughput, normalised)

CANDIDATE	JUDGE SCORE	QUALITY SCORE	TOTAL COST	AVG LATENCY	GLOBAL SCORE	Δ VS LEADER
★ **GPT-5.4-mini**	53.25	60.05	—	2115 ms	**67.28**	—
GPT-5.4	56.25	62.10	—	3820 ms	**54.38**	-12.90

SCORE BY CATEGORY

CATEGORY	GPT-5.4-MINI	GPT-5.4
Coding	99.0	98.0
Reasoning	88.0	95.0
Structured Output	8.0	8.0

CATEGORY	GPT-5.4-MINI	GPT-5.4
Translation	18.0	24.0

VISUAL ANALYSIS

Sub-Score Radar

Dimension-level averages — shows the "personality" of each model.

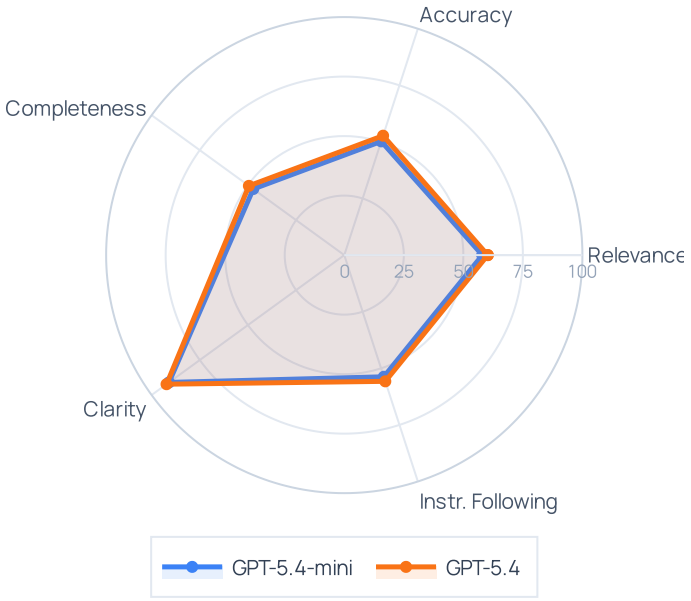

Efficiency Frontier — Latency vs Quality

Upper-left is ideal: high quality at low latency.

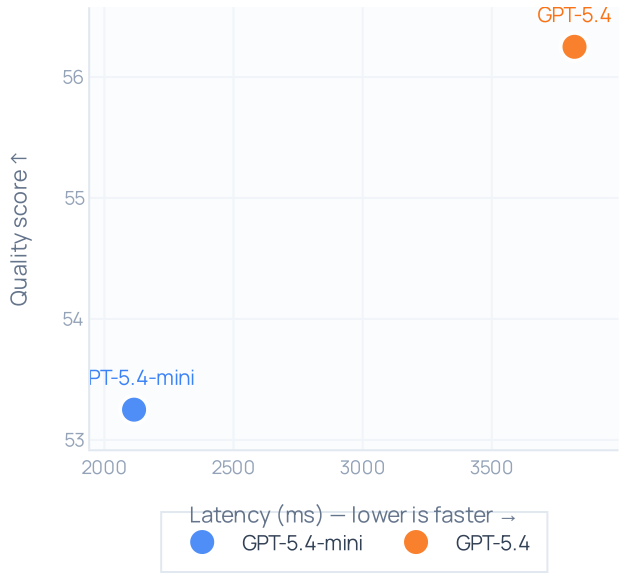

Category Efficiency Radar

Average judge score per prompt category.

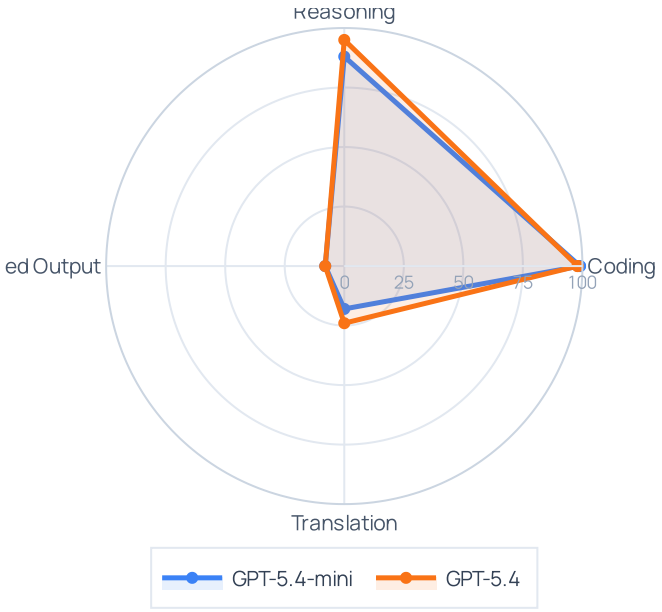

CANDIDATE SECTIONS

GPT-5.4-mini

openai / remote

67.28	2115 ms	51.87 tok/s	—
GLOBAL SCORE	AVG LATENCY	THROUGHPUT	TOTAL COST

DIMENSION	AVG SCORE
Relevance	58.00
Accuracy	50.25
Completeness	47.25
Clarity	91.00
Instruction Following	53.75

Average overall score 53.25 across 4 prompts with mean latency 2115 ms.

STRENGTHS

Clear and understandable response.; Politely requests the missing raw logs needed to perform the task.; Identifies ad hominem/poisoning the well correctly.

WEAKNESSES

Does not attempt the required conversion task.; Does not return any JSON output.; Does not address timestamp normalization, severity mapping, DEBUG omission, or private IP masking.

GPT-5.4

openai / remote

54.38	3820 ms	32.95 tok/s	—
GLOBAL SCORE	AVG LATENCY	THROUGHPUT	TOTAL COST

DIMENSION	AVG SCORE
Relevance	60.25
Accuracy	52.75
Completeness	49.50
Clarity	92.25
Instruction Following	55.75

Average overall score 56.25 across 4 prompts with mean latency 3820 ms.

STRENGTHS

Clear and understandable response.; Politely requests the missing raw logs needed to perform the task.; Directly identifies ad hominem/motive fallacy, matching the benchmark well.

WEAKNESSES

Does not attempt the required conversion task.; Does not return any JSON output.; Does not address timestamp normalization, severity mapping, DEBUG omission, or private IP masking.

PROMPT SECTIONS

transform_messy_logs_json

Structured Output

USER PROMPT

```
Convert the following raw server logs into a valid JSON array. Each object
must have a `timestamp` in ISO 8601, a `severity` level mapped to an
integer (0-5), and a `metadata` object containing extracted IP addresses.
IMPORTANT: If a log line contains 'DEBUG', omit it entirely. If an IP is
from a private range (192.168.x.x), mask it as 'INTERNAL'.
```

Evaluation criteria: Reward strict adherence to the masking and omission rules. Penalize any invalid JSON or missed 'DEBUG' lines.

CANDIDATE	SCORE	RANK	LATENCY	TOKENS	COST	QUAL/100 TOK
GPT-5.4	**8.00**	#1	1846 ms	106	—	7.55
GPT-5.4-mini	**8.00**	#2	1500 ms	106	—	7.55

GPT-5.4

```
Please provide the raw server logs you want converted.
```

Strengths: Clear and understandable response.; Politely requests the missing raw logs needed to perform the task.

Weaknesses: Does not attempt the required conversion task.; Does not return any JSON output.; Does not address timestamp normalization, severity mapping, DEBUG omission, or private IP masking.; Fails to follow the benchmark instruction to process the provided task.

The response is easy to understand, but it does not execute the benchmark task at all. Instead of converting logs into a valid JSON array with ISO 8601 timestamps, mapped severity integers, and extracted/masked IP metadata, it asks the user to provide the raw logs. Since the benchmark evaluates adherence to omission and masking rules and expects an output transformation, this is effectively a non-answer. Its only positive aspect is clarity.

GPT-5.4-mini

```
Please provide the raw server logs you want converted.
```

Strengths: Clear and understandable response.; Politely requests the missing raw logs needed to perform the task.

Weaknesses: Does not attempt the required conversion task.; Does not return any JSON output.; Does not address timestamp normalization, severity mapping, DEBUG omission, or private IP masking.; Fails to follow the benchmark instruction to process the provided task.

The response is easy to understand, but it does not execute the benchmark task at all. Instead of converting logs into a valid JSON array with ISO 8601 timestamps, mapped severity integers, and extracted/masked IP metadata, it asks the user to provide the raw logs. Since the benchmark evaluates adherence to omission and masking rules and expects an output transformation, this is effectively a non-answer. Its only positive aspect is clarity.

critique_fallacious_argument

Reasoning

USER PROMPT

```
Analyze the following argument: 'Our competitors are pushing for AI
regulation because they are scared of our innovation. If we regulate AI
now, we are essentially inviting a digital dark age. Every major tech
leader who supports regulation also has a history of failed product
launches.' Identify at least three logical fallacies used here and explain
why they weaken the argument.
```

Evaluation criteria: Reward correct identification of Ad Hominem, Slippery Slope, and Straw Man/Motivated Reasoning. Penalize vague 'it sounds biased' answers.

CANDIDATE	SCORE	RANK	LATENCY	TOKENS	COST	QUAL/100 TOK
GPT-5.4	**95.00**	#1	8334 ms	611	—	15.55
GPT-5.4-mini	**88.00**	#2	3102 ms	402	—	21.89

GPT-5.4

```
Here are at least three logical fallacies in the argument, along with
why they weaken it:

1. **Ad hominem / Motive fallacy**
   - **Claim:** “Our competitors are pushing for AI regulation because
they are scared of our innovation.”
   - **Why it’s fallacious:** This attacks the competitors’ motives
instead of addressing whether regulation itself might be reasonable.
Even if competitors were motivated by fear or self-interest, that would
not prove regulation is bad.
   - **Why it weakens the argument:** It replaces evidence-based
reasoning with speculation about intent.

2. **Slippery slope**
   - **Claim:** “If we regulate AI now, we are essentially inviting a
digital dark age.”
   - **Why it’s fallacious:** This assumes that regulation will
inevitably lead to an extreme negative outcome without showing the
steps in between. Not all regulation destroys innovation; some
regulation can shape or improve it.
   - **Why it weakens the argument:** It exaggerates consequences
instead of proving them, making the reasoning alarmist rather than
```

```
logical.

3. **Genetic fallacy / Guilt by association**
   - **Claim:** “Every major tech leader who supports regulation also has a history of failed product launches.”
   - **Why it’s fallacious:** This dismisses support for regulation based on who supports it and their past failures, rather than evaluating the merits of their position. Someone can have failed products and still be correct about policy.
   - **Why it weakens the argument:** It avoids engaging with the actual case for regulation.

4. **Hasty generalization**
   - **Claim:** “Every major tech leader who supports regulation also has a history of failed product launches.”
   - **Why it’s fallacious:** This makes a broad claim about “every” major tech leader without evidence. It may rely on selective examples rather than a full, representative sample.
   - **Why it weakens the argument:** Overgeneralized claims are easy to challenge and undermine credibility.

5. **False dilemma**
   - **Implied claim:** Either we avoid regulation or we face a “digital dark age.”
   - **Why it’s fallacious:** It presents only two extreme options, ignoring middle-ground possibilities such as limited, targeted, or adaptive regulation.
   - **Why it weakens the argument:** It oversimplifies a complex policy issue.

Overall, the argument is weak because it relies on attacking people, speculating about motives, and exaggerating consequences instead of providing evidence about whether AI regulation would actually help or harm innovation.
```

Strengths: Directly identifies ad hominem/motive fallacy, matching the benchmark well.; Correctly labels the 'digital dark age' claim as slippery slope.; Explains clearly why personal motives and failed launches do not refute the case for regulation.; Adds useful supporting fallacies like hasty generalization and false dilemma.; Well structured and fully answers the prompt with more than three fallacies.

Weaknesses: Does not explicitly use the term straw man, though motive/genetic fallacy covers part of the intended reasoning.; 'Guilt by association' is slightly less precise than ad hominem/genetic fallacy for the final clause.

Candidate B is the strongest answer because it closely matches the benchmark expectations. It explicitly identifies ad hominem/motive fallacy, correctly labels the 'digital

dark age' claim as a slippery slope, and explains why these moves weaken the argument by replacing evidence with speculation and alarmism. Its treatment of the failed-product-launches line as genetic fallacy/guilt by association is also solid, and the added hasty generalization and false dilemma points improve completeness. While it does not explicitly name straw man, it still captures the motivated-reasoning aspect the benchmark rewards. Overall, it is relevant, clear, and well justified.

GPT-5.4-mini

```
Here are **at least three logical fallacies** in that argument, along
with why they weaken it:

1. **Ad hominem / Poisoning the well**
   - The claim that “every major tech leader who supports regulation
also has a history of failed product launches” attacks the people
making the case rather than addressing the merits of AI regulation
itself.
   - This weakens the argument because a person’s past failures do not
determine whether their current position is correct.

2. **Appeal to fear**
   - Saying that regulating AI would mean “inviting a digital dark age”
is designed to trigger fear rather than provide evidence.
   - This is weak reasoning because it assumes a catastrophic outcome
without showing a direct causal link between regulation and societal
collapse.

3. **Genetic fallacy / Motive fallacy**
   - The statement that competitors want regulation “because they are
scared of our innovation” tries to discredit their position by
speculating about their motives instead of evaluating the regulation
proposal on its own merits.
   - Even if their motives were self-interested, the regulation could
still be justified or beneficial.

4. **Hasty generalization**
   - “Every major tech leader who supports regulation also has a
history of failed product launches” makes a sweeping claim about an
entire group based on an irrelevant pattern.
   - Even if some supporters had failures, that does not prove their
regulatory views are wrong.

Overall, the argument relies more on **emotion, personal attacks, and
speculation about motives** than on evidence or reasoning about the
actual effects of AI regulation.
```

Strengths: Identifies ad hominem/poisoning the well correctly.; Explains clearly why attacking motives or past failures does not address the merits of regulation.; Recognizes the fear-based catastrophic framing in 'digital dark age'.; Provides at least three fallacies and concise explanations.

Weaknesses: Misses the explicitly rewarded slippery slope label, using appeal to fear instead.; Does not identify straw man; motive/genetic fallacy is plausible but less aligned with the evaluation notes.; Hasty generalization point is somewhat redundant with the final sentence about failed product launches.

Candidate A gives a solid analysis with multiple plausible fallacies and good explanations of why they weaken the argument. The ad hominem/poisoning the well point is strong, and the discussion of motive/genetic fallacy is relevant to the claim about competitors being 'scared.' The answer is also clear and well organized. However, under the benchmark notes, the ideal response should identify Ad Hominem, Slippery Slope, and Straw Man/Motivated Reasoning. A labels the 'digital dark age' claim as appeal to fear rather than slippery slope, which is understandable but not the preferred classification here. It also does not explicitly identify straw man. So while the reasoning is mostly correct, it is slightly less aligned with the benchmark target than Candidate B.

python_two_sum

Coding

USER PROMPT

```
Write a Python function `two_sum(nums, target)` that returns the indices of
two numbers whose sum equals `target`. Provide an efficient solution.
```

Evaluation criteria: Reward correct O(n) hashmap-based implementations.

CANDIDATE	SCORE	RANK	LATENCY	TOKENS	COST	QUAL/100 TOK
GPT-5.4-mini	**99.00**	#1	2182 ms	218	—	45.41
GPT-5.4	**98.00**	#2	3419 ms	205	—	47.80

GPT-5.4

```
```python
def two_sum(nums, target):
    seen = {}
    for i, num in enumerate(nums):
        complement = target - num
        if complement in seen:
            return [seen[complement], i]
        seen[num] = i
    return []
```

How it works
- Use a dictionary `seen` to store numbers we've already visited and
their indices.
- For each number, compute the `complement = target - num`.
- If that complement is already in `seen`, we found the pair.
- Return their indices immediately.

Example
```python
print(two_sum([2, 7, 11, 15], 9))  # [0, 1]
```

Time complexity
- **O(n)** time
- **O(n)** space
```

Strengths: Provides a correct O(n) hashmap-based implementation.; Uses the efficient single-pass dictionary solution rewarded by the benchmark notes.; Includes a concise explanation and example.; States correct time and space complexity.

Weaknesses: Slightly less detailed explanation than Candidate A.; Does not explicitly mention the one-pass nature as clearly as A.; Returns an empty list when no solution exists without discussing assumptions.

Candidate B also fully addresses the prompt with the optimal hashmap-based solution. The code is correct, efficient, and follows the expected O(n) approach. It clearly explains the role of the dictionary and includes an example plus complexity notes. Compared with Candidate A, it is marginally less complete in exposition, but the implementation quality is equally strong. The return of an empty list when no pair exists is reasonable, though not explicitly justified.

GPT-5.4-mini

```
```python
def two_sum(nums, target):
    seen = {}  # value -> index

    for i, num in enumerate(nums):
        complement = target - num
        if complement in seen:
            return [seen[complement], i]
        seen[num] = i

    return []
```

How it works
- We scan the list once.
- For each number `num`, we check whether `target - num` was seen before.
- If yes, we return the earlier index and the current index.
- If not, we store the current number and its index in a dictionary.

Efficiency
- **Time complexity:** `O(n)`
- **Space complexity:** `O(n)`

Example
```python
print(two_sum([2, 7, 11, 15], 9))  # [0, 1]
```
```

Strengths: Provides a correct O(n) hashmap-based implementation.; Returns indices as required and handles the lookup-before-insert pattern correctly.; Includes clear explanation, complexity analysis, and example usage.; Matches the prompt exactly and is efficient.

Weaknesses: Returns an empty list when no solution exists without stating whether this behavior is assumed by the prompt.

Candidate A fully satisfies the benchmark prompt. The function uses the standard optimal hashmap approach with a single pass through the list, achieving O(n) time and O(n) space. The implementation is correct, including the important detail of checking for the complement before inserting the current number, which avoids using the same element twice improperly. The explanation is clear and complete, and the response includes both complexity analysis and a working example. The only minor limitation is that it returns an empty list if no solution exists, which is acceptable in many contexts but not explicitly discussed.

translate_fr_to_en_business

Translation

USER PROMPT

```
Translate the following text from French to English using a natural and
professional tone.
```

Evaluation criteria: Reward faithfulness, fluent English, and professional register.

CANDIDATE	SCORE	RANK	LATENCY	TOKENS	COST	QUAL/100 TOK
GPT-5.4	**24.00**	#1	1681 ms	39	—	61.54
GPT-5.4-mini	**18.00**	#2	1676 ms	37	—	48.65

GPT-5.4

```
Of course — please provide the French text you'd like me to translate.
```

Strengths: Clear, polite, and professional wording.; Slightly more direct and concise than Candidate A.

Weaknesses: Did not perform the requested translation.; Still fails to satisfy the core task due to absence of translated output.; Includes conversational filler not needed for benchmark performance.

This response is clear, natural, and professional, and it handles the missing input appropriately for a live conversation. However, as a benchmark answer it does not fulfill the actual instruction, which was to translate French text into English. Because no translation is provided, completeness and accuracy are very low. It ranks above A only because it is marginally more direct and polished.

GPT-5.4-mini

```
Of course — please paste the French text you'd like translated.
```

Strengths: Clear and polite phrasing.; Natural English and professional tone.

Weaknesses: Did not perform the requested translation.; Adds an unnecessary preface ('Of course').; Fails to address the missing-source issue in a benchmark-appropriate way.

The response is fluent and professionally phrased, but it does not translate anything because no French source text was provided. While asking for the missing text is reasonable in an interactive setting, it does not satisfy the benchmark instruction to translate the text. It is therefore minimally relevant, largely incomplete, and weak on instruction following.